



4. Gradient based Learning

- Networks of **continuous units**
- **Regression** problems
- **Gradient descent**, *backpropagation of error*
- The role of the **learning rate**
- **Online learning**, *stochastic approximation*



input layer (external stimulus)



layered architecture
(here: 6-3-4-1)

directed connections
(here: only to next layer)


hidden units
(internal representation)



output unit(s)
(function of input vector)

$$S_i = g \left(\sum w_{ij} S_j \right)$$

↑ previous layer only



input units $\xi_j \in \mathbb{R}, \xi \in \mathbb{R}^N$

$w_j^{(k)}$ input to hidden weights

hidden layer units

$$\sigma_k = g \left(\sum w_j^{(k)} \xi_j \right)$$

v_k hidden to output weights

S single output unit

output = **non-linear** function of input variables:

$$S = g \left(\sum_{k=1}^K v_k \sigma_k \right) = g \left(\sum_{k=1}^K v_k g \left(\sum w_j^{(k)} \xi_j \right) \right)$$

parameterized by set of all weights (and threshold)

continuous activation functions, e.g. $g(x) = \tanh(\gamma x)$
for all nodes in the network

given a network architecture, the weights (and thresholds)
parameterize a continuous function:

$$\xi \in \mathbb{R}^N \rightarrow \sigma(\xi) \in \mathbb{R} \quad (\text{here: single output unit})$$

Learning as **regression problem**

set of examples $\left\{ \xi^\mu, \tau(\xi^\mu) \right\}_{\mu=1}^P$ with cont. labels $\tau \in \mathbb{R}$

training:

(approximately) implement $\sigma(\xi^\mu) = \tau(\xi^\mu)$ for all μ

generalization:

application to novel data $\sigma(\xi) \approx \tau(\xi)$


training strategy: employ an **error measure**
for comparison of student/teacher outputs

just one very popular and plausible choice:

quadratic deviation: $e(\sigma, \tau) = \frac{1}{2} (\sigma - \tau)^2$

cost function: $E = \frac{1}{P} \sum_{\mu=1}^P e^{\mu} = \frac{1}{P} \sum_{\mu=1}^P \frac{1}{2} (\sigma(\xi^{\mu}) - \tau(\xi^{\mu}))^2$

- defined for a given set of example data
- guides the training process
- is a **differentiable function** of weights and thresholds
- training by **gradient descent** minimization of E



$$\xi_j \in \mathbb{R}, \boldsymbol{\xi} \in \mathbb{R}^N$$


$$\mathbf{w} \in \mathbb{R}^N$$

$$\sigma = g \left(\sum_{j=1}^N w_j \xi_j \right)$$

$$E(\mathbf{w}) = \frac{1}{P} \sum_{\mu=1}^P \frac{1}{2} \left(g(\mathbf{w} \cdot \boldsymbol{\xi}^{\mu}) - \tau(\boldsymbol{\xi}^{\mu}) \right)^2$$

$$\frac{\partial E}{\partial w_k} = \frac{1}{P} \sum_{\mu=1}^P \left(g(\mathbf{w} \cdot \boldsymbol{\xi}^{\mu}) - \tau(\boldsymbol{\xi}^{\mu}) \right) g'(\mathbf{w} \cdot \boldsymbol{\xi}^{\mu}) \xi_k^{\mu}$$

$$\nabla_{\mathbf{w}} E = \frac{1}{P} \sum_{\mu=1}^P \left(g(\mathbf{w} \cdot \boldsymbol{\xi}^{\mu}) - \tau(\boldsymbol{\xi}^{\mu}) \right) g'(\mathbf{w} \cdot \boldsymbol{\xi}^{\mu}) \boldsymbol{\xi}^{\mu}$$



$$\xi_j \in \mathbb{R}, \boldsymbol{\xi} \in \mathbb{R}^N$$

$$\mathbf{w} \in \mathbb{R}^N$$

$$\sigma = g \left(\sum_{j=1}^N w_j \xi_j \right)$$

frequent choice: $g(\mathbf{x}) = \tanh(\mathbf{x})$ $g'(\mathbf{x}) = 1 - \tanh^2(\mathbf{x})$

$$\nabla_{\mathbf{w}} E = \frac{1}{P} \sum_{\mu=1}^P \left(\tanh(\mathbf{w} \cdot \boldsymbol{\xi}^{\mu}) - \tau(\boldsymbol{\xi}^{\mu}) \right) \left(1 - \tanh^2(\mathbf{w} \cdot \boldsymbol{\xi}^{\mu}) \right) \boldsymbol{\xi}^{\mu}$$


single example contribution:

$$\nabla_{\mathbf{w}} e^{\mu} = \left(\tanh(\mathbf{w} \cdot \boldsymbol{\xi}^{\mu}) - \tau(\boldsymbol{\xi}^{\mu}) \right) \left(1 - \tanh^2(\mathbf{w} \cdot \boldsymbol{\xi}^{\mu}) \right) \boldsymbol{\xi}^{\mu}$$

Backpropagation of Error

convenient calculation of the gradient in multilayer networks (\leftarrow chain rule)

example: continuous two-layer network with K hidden units



Exercise: derive $\nabla_{\mathbf{w}_k} E$ and $\frac{\partial E}{\partial v_k}$

the weights \mathbf{w}_k and v_k are used ...

- *downward* for the calculation of hidden states and output
- *upward* for the calculation of the gradient


negative gradient gives the **direction of steepest descent** in E

simple gradient based minimization of E :

sequence $\mathbf{w}_0 \rightarrow \mathbf{w}_1 \rightarrow \dots \rightarrow \mathbf{w}_t \rightarrow \mathbf{w}_{t+1} \rightarrow \dots$

with $\mathbf{w}_{t+1} = \mathbf{w}_t - \eta \nabla E|_{\mathbf{w}_t}$

approaches some minimum of E (?)



learning rate rate η

- controls the step size of the algorithm
- has to be small enough to ensure convergence
- should be as large as possible to facilitate fast learning

assume E has a (local) minimum in \mathbf{w}^* , Taylor expansion in the vicinity:

$$E(\mathbf{w}) \approx E(\mathbf{w}^*) + (\mathbf{w} - \mathbf{w}^*)^T \underbrace{\nabla E|_{*}}_{=0} + \frac{1}{2} (\mathbf{w} - \mathbf{w}^*)^T H^* (\mathbf{w} - \mathbf{w}^*) + \dots$$

$$E(\mathbf{w}) \approx E(\mathbf{w}^*) + \frac{1}{2} (\mathbf{w} - \mathbf{w}^*)^T H^* (\mathbf{w} - \mathbf{w}^*) \quad \nabla E|_{\mathbf{w}} \approx H^* (\mathbf{w} - \mathbf{w}^*)$$

with the positive definite **Hesse matrix** of second derivatives $H^*_{ij} = \left. \frac{\partial^2 E}{\partial w_i \partial w_j} \right|_*$

H^* has only pos. eigenvalues $\lambda_i > 0$, orthonormal eigenvectors \mathbf{u}_i (all $\lambda_i \leq \lambda_{max}$)

gradient descent in the vicinity of \mathbf{w}^* :

$$\mathbf{w}_t - \mathbf{w}^* \equiv \boldsymbol{\delta}_t = \boldsymbol{\delta}_{t-1} - \eta \nabla E|_{\mathbf{w}_{t-1}}$$

$$\boldsymbol{\delta}_t \approx [I - \eta H^*] \boldsymbol{\delta}_{t-1} \approx [I - \eta H^*]^t \boldsymbol{\delta}_0$$

$$\text{expansion in } \{\mathbf{u}_i\}: \boldsymbol{\delta}_0 = \sum_i a_i \mathbf{u}_i$$

$$\boldsymbol{\delta}_t \approx \sum_i a_i [I - \eta H^*]^t \mathbf{u}_i = \sum_i a_i [1 - \eta \lambda_i]^t \mathbf{u}_i$$

with $\mathbf{u}_j^T \mathbf{u}_k = \delta_{jk}$ we obtain

$$|\boldsymbol{\delta}_t|^2 = \sum_i a_i^2 [1 - \eta \lambda_i]^{2t}$$

in detail:

$$w_t = w_{t-1} - \eta \nabla E|_{w_{t-1}}$$

$$\delta_t = \delta_{t-1} - \eta \nabla E|_{w_{t-1}} \approx [I - \eta H^*] \delta_{t-1} \approx [I - \eta H^*]^t \delta_o$$

$$\delta_o = \sum_i a_i \mathbf{u}_i \quad H^* \mathbf{u}_i = \lambda_i \mathbf{u}_i \quad (1 - \eta H^*) \mathbf{u}_i = (1 - \eta \lambda_i) \mathbf{u}_i$$

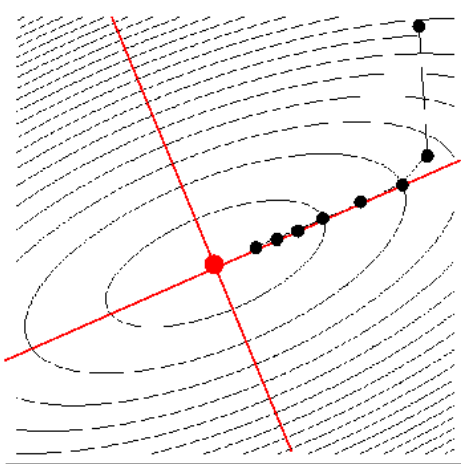
$$\delta_t \approx [I - \eta H^*]^t \sum_i a_i \mathbf{u}_i = \sum_i a_i (1 - \eta \lambda_i)^t \mathbf{u}_i$$

$$\begin{aligned} \delta_t^2 = \delta_t \cdot \delta_t &= \sum_{i,j} a_i [1 - \eta \lambda_i]^t a_j [1 - \eta \lambda_j]^t \underbrace{\mathbf{u}_i \cdot \mathbf{u}_j}_{\substack{=1 \text{ if } i=j \\ =0 \text{ else}}} \\ &= \sum_i a_i^2 [1 - \eta \lambda_i]^{2t} \end{aligned}$$

iteration approaches the minimum, $\lim_{t \rightarrow \infty} |\delta_t| = 0$, only if $|1 - \eta\lambda_i| < 1$ for all i

condition for (local) convergence: $\eta < \eta_{max} = \frac{2}{\lambda_{max}}$ (largest eigenvalue of H^*)


$$\eta < \frac{\eta_{max}}{2} = \frac{1}{\lambda_{max}}$$



$$1 - \eta\lambda_{max} > 0$$

smooth convergence


$$\frac{1}{\lambda_{max}} < \eta < \frac{2}{\lambda_{max}}$$



$$1 - \eta\lambda_{max} < 0$$

oscillations

$$\eta > \frac{2}{\lambda_{max}}$$



$$1 - \eta\lambda_{max} < -1$$

divergence

... the above considerations

- are only valid close to the minimum
local minima can have completely different characteristics (λ_{max})
- do not concern global convergence properties
e.g. the choice of the learning rate far from a minimum

potential problems:

- E can have (many) local minima far from global optimality
- initial conditions determine which minimum will be approached
- anisotropic curvatures can cause strong oscillations
- E can have *saddle points* with $\nabla E = 0$ and/or *flat regions* with $\nabla E \approx 0$
gradient learning can slow down drastically by, e.g., *plateau states*, see below

some modifications:

- improved gradient descent: e.g. time dependent $\eta(t)$
momentum: $\Delta \mathbf{w}_{t+1} = -\eta \nabla E + a \Delta \mathbf{w}_t$ “keep going”
- sophisticated **optimization methods**:
line search procedures, conjugate gradient, second order methods,
e.g. Newton’s method (“matrix update” employs H), ...
- **different learning rates** for different weights, examples:
 - heuristics: $\eta \propto 1/N$ for input-to-hidden, $\eta \propto 1/K$ for hidden-to-output weights
 - simplified version of “matrix update” (assume H is approximately diagonal):
update each weight w_j with a learning rate $\eta_j \propto 1 / \frac{\partial^2 E}{\partial w_j^2}$
 - learning algorithms realize *descent* in E as long as $\Delta \mathbf{w} \cdot \nabla E < 0$
- construction of alternative **well-behaved cost functions**, one example:

$$E = \sum_{\mu} \begin{cases} \gamma (\sigma - \tau)^2 & \text{if } \text{sign}(\sigma) = \text{sign}(\tau) \\ (\sigma - \tau)^2 & \text{if } \text{sign}(\sigma) \neq \text{sign}(\tau) \end{cases} \quad \text{with } \gamma \text{ increasing from 0 to 1.}$$

small γ : emphasis on correct sign of the output large γ : fine tuning of σ

stochastic approximation (on-line gradient descent)

cost function $E = \frac{1}{P} \sum_{\mu=1}^P e^{\mu} \equiv \overline{e^{\mu}}$ is an **empirical average** over examples

→ simple approximation of ∇E by ∇e^{μ} for one example only

- select one $\mu \in \{1, 2, \dots, P\}$ with equal probability $1/P$
- single step: $\mathbf{w}_{t+1} = \mathbf{w}_t + \Delta \mathbf{w}_t = \mathbf{w}_t - \eta \nabla e^{\mu}|_{\mathbf{w}_t}$

- computationally cheap compared to *off-line (batch)* gradient descent
- *intrinsic noise*, fewer problems with local minima, flat regions etc.

(when) does the procedure converge?

behavior close to a (local) minimum \mathbf{w}^* of E ?

averaged learning step:
$$\overline{\Delta \mathbf{w}} = -\eta \overline{\nabla e^\mu|_{\mathbf{w}}} = -\frac{\eta}{P} \sum_{\mu=1}^P \nabla e^\mu|_{\mathbf{w}} = -\eta \nabla E|_{\mathbf{w}}$$

$$\overline{\Delta \mathbf{w}} = 0 \quad \text{for } \mathbf{w} \rightarrow \mathbf{w}^*$$

averaged length of $\Delta \mathbf{w}$:
$$\overline{(\Delta \mathbf{w})^2} = \eta^2 \overline{(\nabla e^\mu|_*)^2} > 0$$


~~€ΔΔΔΔΔ~~ Possible if all $e^\mu = 0$

generic behavior

for constant rate $\eta > 0$:

$$\lim_{t \rightarrow \infty} (\Delta \mathbf{w}_t)^2 > 0$$

(fluctuations remain non-zero)



convergence in the sense of $(\Delta \mathbf{w})^2 \rightarrow 0$ only if $\eta(t) \rightarrow 0$ for $t \rightarrow \infty$

one can show: $\lim_{t \rightarrow \infty} \sum_t \eta(t) \rightarrow \infty$ but $\lim_{t \rightarrow \infty} \sum_t \eta(t)^2 < \infty$ is required

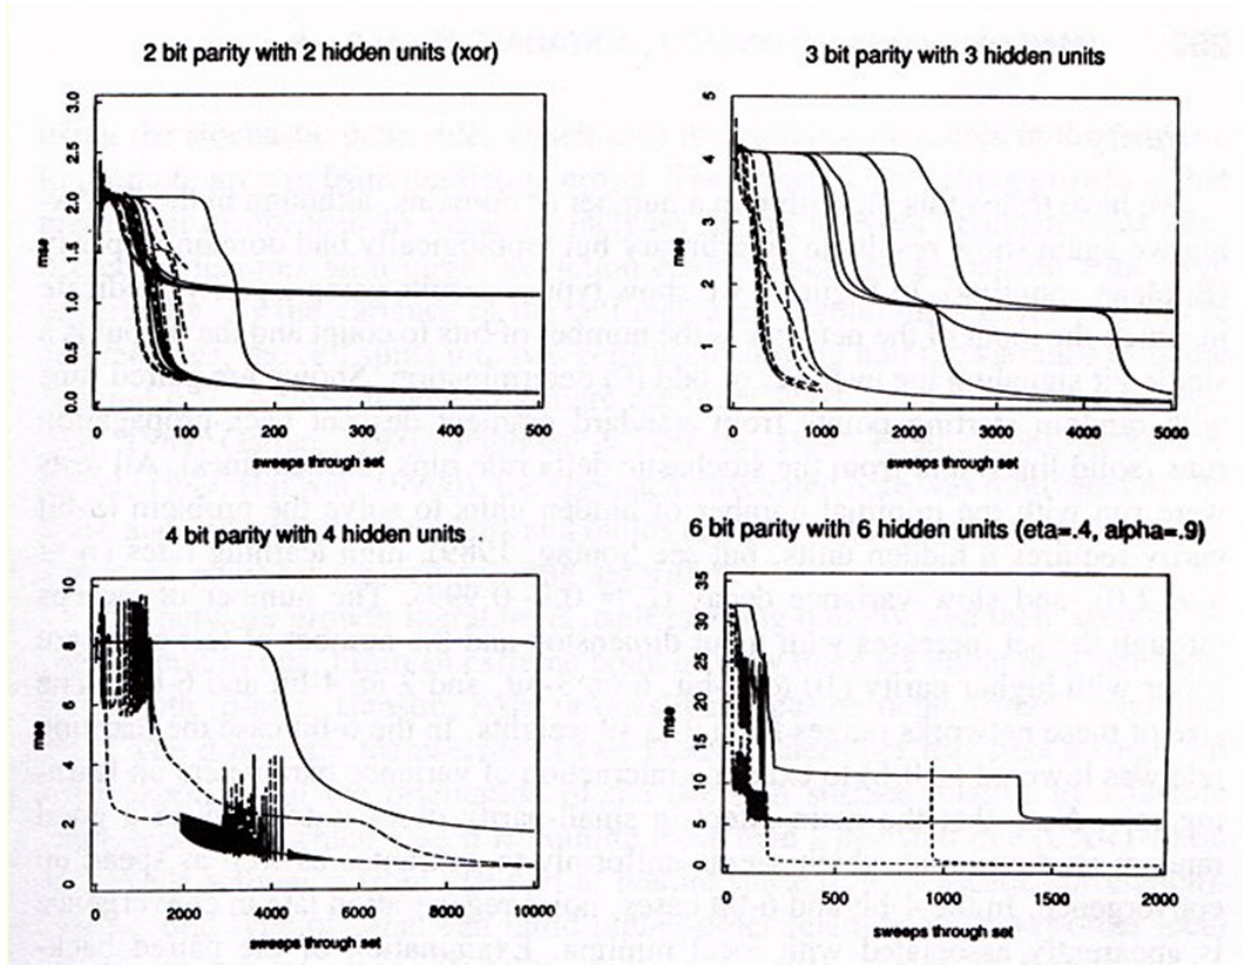
satisfied by, e.g. $\eta(t) \propto \frac{1}{t}$ for large t **learning rate schedules**, e.g. $\eta(t) = \frac{a}{b+t}$

alternative: retain finite learning rate and fluctuations, average weights over time

Plateau states


frequent observation:

training of multilayer networks is delayed by *quasi-stationary plateaus*



(S.J. Hanson, in: Y. Chauvin and D. Rumelhart, *The Handbook of Backpropagation*, 1995)

example: a two-layer network trained from reliable, perfectly realizable data
by on-line gradient descent



number of examples $P/(KN)$

- fast initial decrease of ε_g
- fast asymptotic decrease of $\varepsilon_g \rightarrow 0$
(here: matching complexity)
- plateau state:
unspecialized h.u. with $\mathbf{w}_k \sim \mathbf{w}_o + noise$
have all obtained some (the same)
information about the unknown rule

occurrence of plateaus relates to symmetries:

the network output is invariant under **permutations of hidden units**

perfectly symmetric state corresponds to a flat region (saddle) in E

successful learning requires **specialization** and can be delayed significantly

math. analysis: D. Saad and S. Solla (1995), M. Biehl, P. Riegler, C. Wöhler (1996)

Remarks

- the extension to several output units $\{\sigma^l\}_{l=1}^L$ is non-trivial

the choice of cost function $E = \frac{1}{2} \sum_{l=1}^L (\sigma^l - \tau^l)^2$ seems plausible

but a generalized metric $E = \frac{1}{2} \sum_{k,l} (\sigma^k - \tau^k) A_{kl} (\sigma^l - \tau^l)$

with a suitable $(L \times L)$ -matrix A might be more appropriate

- approximation of a classifier by regression training:

replace $S^\mu = \pm 1$ by $\sigma^\mu = \pm R$ $|R| < 1$,

take $S = \text{sign}(\sigma^\mu)$ as classification output after training

Note:

- the decrease of ε_g can be slower than *necessary*
- for very steep activation functions or $|R| \approx 1$, the gradient becomes non-informative as $\nabla E \approx 0$